

## ✓ Individual Assignment I: Machine Learning Foundation

### Data Preparation

GitHub Repository: <https://github.com/christophrrrr/ML-fundamentals-2026>

This notebook executes data preparation and feature engineering tasks for the UCI Bank Marketing Dataset (`bank-additional.csv`), adhering to data leakage prevention principles.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score, ConfusionMatrixDisplay, classification_report, f1_score
from sklearn.feature_selection import VarianceThreshold

import warnings
warnings.filterwarnings('ignore')

from imblearn.over_sampling import RandomOverSampler
```

## 1. Identifying the Prediction Target

Lecture material: Lecture 1 (Problem Formulation), Lecture 2 (Data Inspection).

**Target Selection:** The target variable is `y`. `y` records "yes" or "no" indicating whether the client subscribed to a term deposit. This aligns with the stated objective of the direct marketing campaigns.

**Invalid Alternatives:** Three other variables might appear to be valid targets but must not be used:

1. `duration`: This represents the call duration in seconds. While highly correlated with `y`, it is an outcome of the call. At prediction time (before or during the start of the call), this information is unavailable. Predicting `duration` does not address the business goal of identifying who will subscribe. Including it results in data leakage.
2. `poutcome`: Records the result of the previous marketing campaign. The prediction objective defined for this assignment is whether the client subscribes in the current campaign. `poutcome` describes an event that has already occurred and is available as an input at prediction time — it is an input feature, not the target. Using it as a prediction target would mean predicting something already known, which has no operational value.
3. `campaign`: Records the number of contacts performed during the current campaign. One might argue this represents campaign effort worth predicting or optimizing. However, `campaign` is a campaign execution variable accumulated during the contact process — it is not the business outcome. The goal is predicting client behavior (`y`), not the number of calls made. It is also partially available at prediction time (current call count), making its use as a target conceptually incoherent.

## ✓ 2. Data Loading and Exploration

Lecture material: Lecture 1 (Problem Formulation), Lecture 2 (Data Inspection and EDA).

- `bank-additional.csv` is the 10% sample (4119 rows) randomly selected from `bank-additional-full.csv` (41188 rows).
- We prefer the full dataset but fall back to the 10% sample to keep computation light or if the full set is unavailable. The preprocessing pipeline remains structurally identical regardless. Note that UCI bank datasets commonly use the semicolon `;` separator.

```
# Load dataset
import os

full_filepath = 'data/bank-additional-full.csv'
sample_filepath = 'data/bank-additional.csv'
github_sample_url = 'https://raw.githubusercontent.com/christophrrrr/ML-fundamentals-2026/main/data/bank-additional.csv'

# Attempt to load full dataset first, then fall back to sample, then to remote link
if os.path.exists(full_filepath):
    print(f"Loading full dataset from: {full_filepath}")
    df = pd.read_csv(full_filepath, sep=';')
elif os.path.exists(sample_filepath):
    print(f"Loading 10% sample dataset from: {sample_filepath}")
    df = pd.read_csv(sample_filepath, sep=';')
else:
    print(f"Local instance not found. Downloading 10% sample directly from GitHub repository...")
    df = pd.read_csv(github_sample_url, sep=';')

# Basic structure
print(f"Number of observations: {df.shape[0]}")
print(f"Number of features: {df.shape[1]}")

print("\n--- Data Types ---")
print(df.dtypes)

print("\n--- Summary Statistics ---")
display(df.describe())
```

```
Local instance not found. Downloading 10% sample directly from GitHub repository...
Number of observations: 4119
Number of features: 21
```

```
-- Data Types --
age           int64
job            object
marital        object
education      object
default        object
housing        object
loan           object
contact        object
month          object
day_of_week    object
duration       int64
campaign       int64
pdays          int64
previous       int64
poutcome       object
emp.var.rate   float64
cons.price.idx float64
cons.conf.idx  float64
euribor3m     float64
nr.employed   float64
y              object
dtype: object
```

```
-- Summary Statistics --
```

	age	duration	campaign	pdays	previous	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	nr.employed	grid
<b>count</b>	4119.000000	4119.000000	4119.000000	4119.000000	4119.000000	4119.000000	4119.000000	4119.000000	4119.000000	4119.000000	
<b>mean</b>	40.113620	256.788055	2.537266	960.422190	0.190337	0.084972	93.579704	-40.499102	3.621356	5166.481695	
<b>std</b>	10.313362	254.703736	2.568159	191.922786	0.541788	1.563114	0.579349	4.594578	1.733591	73.667904	
<b>min</b>	18.000000	0.000000	1.000000	0.000000	0.000000	-3.400000	92.201000	-50.800000	0.635000	4963.600000	
<b>25%</b>	32.000000	103.000000	1.000000	999.000000	0.000000	-1.800000	93.075000	-42.700000	1.334000	5099.100000	
<b>50%</b>	38.000000	181.000000	2.000000	999.000000	0.000000	1.100000	93.749000	-41.800000	4.857000	5191.000000	
<b>75%</b>	47.000000	317.000000	3.000000	999.000000	0.000000	1.400000	93.994000	-36.400000	4.961000	5228.100000	
<b>max</b>	88.000000	3643.000000	35.000000	999.000000	6.000000	1.400000	94.767000	-26.900000	5.045000	5228.100000	

#### Variable Identification (raw dataset, prior to feature engineering):

- Numerical:** age, duration, campaign, pdays, previous, emp.var.rate, cons.price.idx, cons.conf.idx, euribor3m, nr.employed
- Categorical:** job, marital, education, default, housing, loan, contact, month, day\_of\_week, poutcome, y

Note: duration is identified here for completeness but is dropped immediately in the next section due to target leakage. pdays is replaced by two engineered features (prev\_contacted, pdays\_clean) that separate its binary and continuous information.

```
# Target Distribution
y_counts = df['y'].value_counts()
y_pct = df['y'].value_counts(normalize=True)

# --- Figure 1: Target, Age, Job ---
fig, axes = plt.subplots(1, 3, figsize=(16, 4))
df['y'].value_counts().plot(kind='bar', color=['#1f77b4', '#ff7f0e'], ax=axes[0])
axes[0].set_title('Target Variable (y) Distribution')
axes[0].set_ylabel('Count')

df['age'].plot(kind='hist', bins=20, color='skyblue', edgecolor='black', ax=axes[1])
axes[1].set_title('Age Distribution')

df['job'].value_counts().plot(kind='bar', color='lightgreen', edgecolor='black', ax=axes[2])
axes[2].set_title('Job Category Distribution')
plt.tight_layout()
plt.show()

# --- Figure 2: Campaign, Previous, Education, Marital ---
fig2, axes2 = plt.subplots(2, 2, figsize=(16, 10))

df['campaign'].plot(kind='hist', bins=30, color='salmon', edgecolor='black', ax=axes2[0, 0])
axes2[0, 0].set_title('Campaign (Number of Contacts)')

df['previous'].plot(kind='hist', bins=10, color='violet', edgecolor='black', ax=axes2[0, 1])
axes2[0, 1].set_title('Previous Contacts')

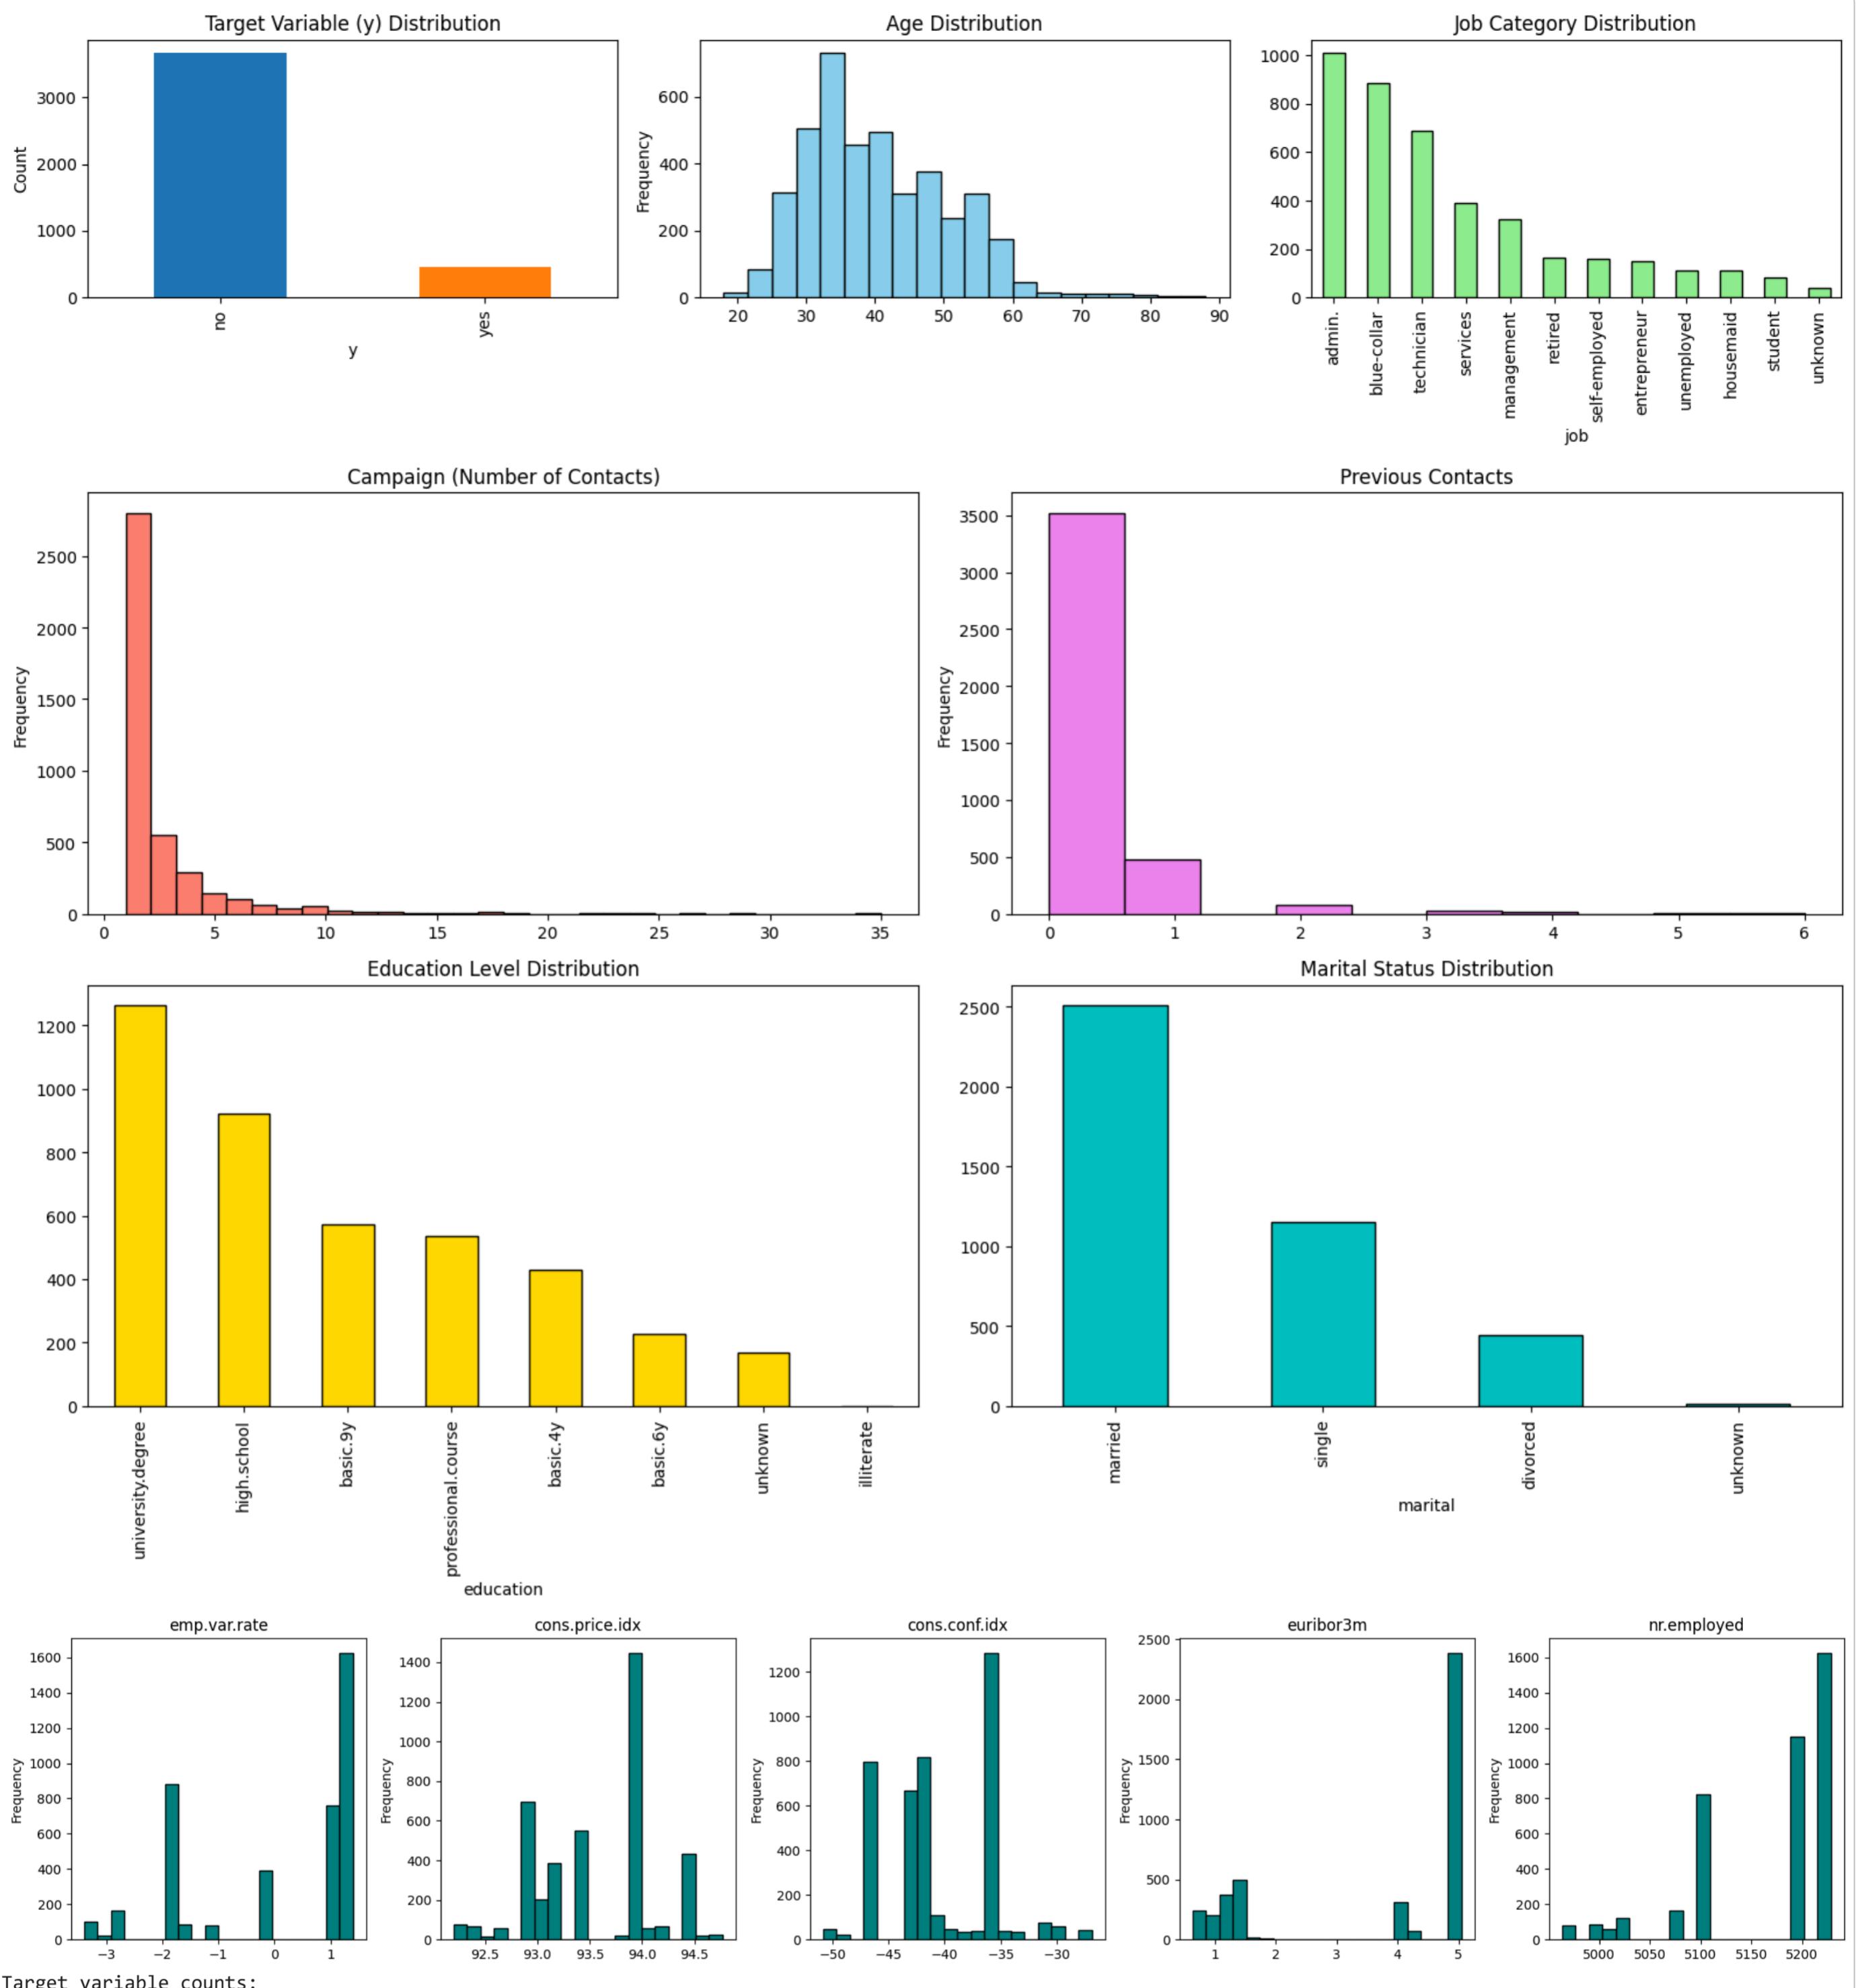
df['education'].value_counts().plot(kind='bar', color='gold', edgecolor='black', ax=axes2[1, 0])
axes2[1, 0].set_title('Education Level Distribution')

df['marital'].value_counts().plot(kind='bar', color='c', edgecolor='black', ax=axes2[1, 1])
axes2[1, 1].set_title('Marital Status Distribution')

plt.tight_layout()
plt.show()

# --- Figure 3: Macroeconomic Variables ---
macro_vars = ['emp.var.rate', 'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed']
fig3, axes3 = plt.subplots(1, 5, figsize=(20, 4))
for i, var in enumerate(macro_vars):
    df[var].plot(kind='hist', bins=20, color='teal', edgecolor='black', ax=axes3[i])
    axes3[i].set_title(var)
plt.tight_layout()
plt.show()

print("Target variable counts:")
print(y_counts)
print("\nTarget variable percentages:")
print(y_pct)
```



#### Target variable counts:

```
y
no    3668
yes   451
Name: count, dtype: int64
```

#### Target variable percentages:

```
y
no    0.890507
yes   0.109493
Name: proportion, dtype: float64
```

#### General Observations:

- Class Imbalance:** Only ~10.9% of clients subscribed (`yes`). Class imbalance handling is required to prevent the model from trivializing predictions.
- Skewed Variables:** `campaign` is right-skewed (most clients are contacted 1–3 times, with a long tail reaching 35). `previous` is zero for the majority of clients. A log transform of `campaign` was considered to reduce the influence of outliers on the linear model's coefficient estimation. It was ultimately not applied here to keep the pipeline interpretable and because StandardScaler will centre and scale the variable regardless — but this remains a worthwhile preprocessing step to explore.
- Category Ratios:** `university.degree` and `high.school` represent the most frequent `education` levels. The majority of clients are married.
- Target Leakage Variable:** `duration` is only known after the call finishes. It must be dropped.
- Implicit Missing Values:** Categorical variables utilize "unknown" as an implicit missing value. `pdays` uses 999 to indicate "never contacted before".

## Macroeconomic Observations:

- `euribor3m` and `nr.employed` show bimodal distributions, clustering into two distinct regimes that likely correspond to pre- and post-2008 economic periods in the dataset.
- `emp.var.rate` is similarly clustered rather than continuous, reinforcing that macroeconomic features track the same underlying economic cycle.
- These distributions suggest the macroeconomic block may carry redundant information — addressed formally in Feature Selection.

## 3. Task Ordering

Lecture material: Lecture 2 (Data Splitting and Leakage), Lecture 5 (Preprocessing), Lecture 9 (ML Pipeline).

To prevent data leakage, data preparation tasks are executed in the following sequence:

### 1. Identifying Target & Data Loading (Completed above)

- Allowed: Raw dataset viewing, broad target assessment.
- Not allowed: Predicting on or analyzing combinations of target vs features globally.
- Leakage risk if changed: None at this stage, assuming `duration` is dropped manually before modeling metrics run.

### 2. Managing Missing Values (Identification & Structural Cleaning)

- Allowed: Finding distinct string literals (e.g. "unknown") or sentinel values (999) and structurally replacing them with `NaN` or indicator flags.
- Not allowed: Computing median, mean, or mode across the column to fill the `NaN` values.
- Leakage risk if changed: Replacing "unknown" specifically does not use global distribution data. However, if statistical imputation were performed here instead, it would leak test-set central tendencies into the training data.

### 3. Data Splitting

- Allowed: Raw input variables (`X`) and targets (`y`).
- Not allowed: Any fitted statistical boundaries, encodings, or synthetic samples.
- Leakage risk if changed: If delayed, transformation steps would consume information belonging to the test set, compromising final evaluation integrity.

### 4. Managing Missing Values (Statistical Imputation)

- Allowed: Medians/Modes calculated from `X_train`.
- Not allowed: Test set distribution properties.
- Leakage risk if changed: If placed before Data Splitting, the median would include test observations.

### 5. Encoding Categorical Variables

- Allowed: List of distinct categories present in `X_train`.
- Not allowed: Categories that only exist in `X_test`.
- Leakage risk if changed: The algorithm would map dummy dimensions for categories it has not seen yet.

### 6. Feature Scaling

- Allowed: Compute mean and variance only over `X_train` via `.fit()`.
- Not allowed: Running `.fit()` on `X_test`.
- Leakage risk if changed: If placed before Data Splitting, the feature distances for the test set observations would be compressed based on training outliers.

### 7. Feature Selection

- Allowed: Variance thresholds and correlation matrices computed over `X_train`.
- Not allowed: Entire dataset correlations.
- Leakage risk if changed: If executed upfront, variables would be deleted based on how they correlate with target labels inside the test set.

### 8. Addressing Class Imbalance

- Allowed: Resampling methods (e.g. RandomOverSampler) applied within the training set.
- Not allowed: Resampling before Data Splitting.
- Leakage risk if changed: If RandomOverSampler generated synthetic samples before the train/test split, synthesized points mathematically linked to train observations would land in the test set.

**Incorrect Ordering Example (Scaling before Splitting):** If Feature Scaling is performed before Data Splitting, the mean and standard deviation are calculated across the entire dataset. The standardized test values inherently contain information about the central tendency of the training set. This is data leakage.

**Incorrect Ordering Example 2 — Resampling before splitting:** If RandomOverSampler were applied to the full dataset before the train/val/test split, synthetic minority samples would be generated using the entire data distribution. When the dataset is subsequently split, some synthetic samples — constructed using observations that end up in the validation or test sets — will appear in the training set. The model trains on data derived from the test set. Validation and test metrics will be inflated because the boundary between training and evaluation data has been contaminated. The correct position for any resampling operation is after splitting, applied to the training set only.

## 4. Managing Missing Values (Part 1: Identification & Sentinel Cleaning)

Lecture material: Lecture 2 (Data Inspection), Lecture 5 (Preprocessing and Pipeline Discipline).

### Identification:

- Explicit missing values (`NaN`) are largely absent in this CSV.
- Implicit missing values are abundant. Words like "unknown" map strictly to missing information.
- In numerical columns, `pdays=999` acts as a sentinel for "client was not previously contacted".

We must convert these implicit symbols into standard structural missingness (`NaN`) before splitting, alongside creating feature flags.

**Note on Leakage:** Because we are just structurally replacing "unknown" → `NaN` and extracting `pdays != 999`, we are not calculating statistics. Therefore, this is purely "data cleaning" and is safe to execute before Data Splitting.

```

# Drop 'duration' immediately to prevent leakage before any further inspection
if 'duration' in df.columns:
    df = df.drop(columns=['duration'])

# Count categorical 'unknown' and numerical '999' before cleaning
missing_counts = []

for col in df.drop(columns=['y']).select_dtypes(include=['object']).columns:
    unknown_count = (df[col] == 'unknown').sum()
    if unknown_count > 0:
        missing_counts.append({
            'Variable': col,
            'Implicit Missing': unknown_count,
            '% of Total': f"{{(unknown_count / len(df)) * 100:.2f}}%"})
    })

pdays_count = (df['pdays'] == 999).sum()
if pdays_count > 0:
    missing_counts.append({
        'Variable': 'pdays',
        'Implicit Missing': pdays_count,
        '% of Total': f"{{(pdays_count / len(df)) * 100:.2f}}%"})

missing_df = pd.DataFrame(missing_counts)
print("--- Implicit Missing Values Summary Before Structural Cleaning ---")
display(missing_df)

```

--- Implicit Missing Values Summary Before Structural Cleaning ---

	Variable	Implicit Missing	% of Total	Actions
0	job	39	0.95%	
1	marital	11	0.27%	
2	education	167	4.05%	
3	default	803	19.50%	
4	housing	105	2.55%	
5	loan	105	2.55%	
6	pdays	3959	96.12%	

Nächste Schritte: [New interactive sheet](#)

```

# 1. Handle Categorical 'unknown'
cat_cols = df.select_dtypes(include=['object']).columns
for col in cat_cols:
    df[col] = df[col].replace('unknown', np.nan)

# 2. Handle 'pdays' Sentinel
# Create a logical flag for previous contact
df['prev_contacted'] = (df['pdays'] != 999).astype(int)
# Clean the magnitude (so 999 doesn't distort linear models)
df['pdays_clean'] = df['pdays'].replace(999, np.nan)
df = df.drop(columns=['pdays'])

missing_summary = df.isna().sum()
print("Missing (NaN) counts after structured cleaning:")
print(missing_summary[missing_summary > 0])

```

```

Missing (NaN) counts after structured cleaning:
job          39
marital       11
education     167
default       803
housing       105
loan          105
pdays_clean   3959
dtype: int64

```

## 5. Data Splitting

Lecture material: Lecture 2 (Data Splitting and Leakage), Lecture 9 (ML Pipeline).

The independent features and the target are separated using a stratified split.

### Proportions:

- Training: 70% (Used to learn parameters for imputation, scaling, encoding, and modeling).
- Validation: 15% (Used to evaluate model health during iterations and tune hyperparameters).
- Test: 15% (Held-out subset for final generalization reporting).

**Stratification:** `stratify=y` is used because the target is imbalanced (~11% positives). A random split could yield a training set with very few positive examples, leading to instability.

**Leakage Prevention:** Executing this split here ensures that upcoming steps (Scaling, Imputation, Encoding) can only `fit()` on mathematical properties present in `X_train`.

```

X = df.drop(columns=['y'])
y = df['y']

# First split: Train (70%), Temp (30%)
X_train, X_temp, y_train, y_temp = train_test_split(
    X, y, test_size=0.30, stratify=y, random_state=42
)

```

```

# Second split: Temp -> Validation (50% of 30% = 15%) and Test (15%)
X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.50, stratify=y_temp, random_state=42
)

print(f"X_train shape: {X_train.shape}")
print(f"X_val shape: {X_val.shape}")
print(f"X_test shape: {X_test.shape}")

X_train shape: (2883, 20)
X_val shape: (618, 20)
X_test shape: (618, 20)

```

## 6. Managing Missing Values (Part 2: Imputation)

The statistical modeling logic is established inside `scikit-learn` Pipelines.

- **Numerical Imputation (`pdays_clean`):** `pdays_clean` has 96% missing data. Because the variance of "was previously contacted vs wasn't" is captured using the binary `prev_contacted` flag, imputing `pdays_clean` with the Train median imputes a near-empty column with a static baseline. The variable is retained because the 4% of clients with previous campaigns possess numerical magnitudes that a linear model can use.
- **Categorical Imputation:** We replace categorical `NaN` with the explicit string `"missing"`. This records missingness directly as an additional feature state rather than discarding it.
- **Note on `default` (19.5% missing):** `default` records whether a client has credit in default — the highest missingness rate of any categorical variable. Dropping the column entirely was considered, but rejected: the 80.5% of non-missing values still carry signal, and the `"missing"` imputation strategy preserves the missingness pattern as its own category, which the model can learn from. If missingness on `default` is itself predictive of subscription behaviour (e.g. banks may withhold this information for certain client profiles), encoding it explicitly as `"missing"` captures that signal rather than discarding it.

```

num_cols = X_train.select_dtypes(include=[np.number]).columns.tolist()
cat_cols = X_train.select_dtypes(exclude=[np.number]).columns.tolist()

# Define Imputers
num_imputer = SimpleImputer(strategy='median')
cat_imputer = SimpleImputer(strategy='constant', fill_value='missing')

# We hold off assembling the full ColumnTransformer until we define Scaling-Encoding.

```

## 7. Encoding Categorical Variables

*Lecture material: Lecture 4 (Categorical Encoding), Lecture 6 (Linear Models).*

**Classification:**

- **Nominal Variables** (e.g., `job`, `marital`, `contact`, `month`): No intrinsic mathematical order.
- **Ordinal Variables** (e.g., `education`): Intrinsic order (`basic.4y` < `high.school` < `university.degree`).

**Strategy — Education (Ordinal vs One-Hot):** `education` is a genuine ordinal variable. `OrdinalEncoder` would assign integer ranks `[0, 1, 2, ... 7]`, which assumes uniform interval spacing between consecutive levels. There is no empirical basis for claiming that the educational jump from `basic.4y` to `basic.6y` carries the same predictive weight as the jump from `high.school` to `university.degree`. Because Logistic Regression relies on linear combinations of its inputs, an incorrectly-spaced ordinal encoding would systematically distort the coefficient for `education`, biasing all predictions that depend on it.

**One-Hot Encoding** is therefore applied to `education` — and all other categorical variables — as a non-parametric alternative. This gives the model a free, independent coefficient for every category, making no assumptions about inter-level distances. The cost is additional dimensionality, but for a dataset of this size (~4000 observations, ~60 post-encoding features) this is well within the capacity of Logistic Regression with L2 regularization.

*Impact on Dimensionality:* Expands categorical columns from 10 variables to ~51 binary features. *Impact on Interpretability:* The Logistic Regression yields a discrete coefficient for each category (e.g., `job_retired`, `education_university.degree`), making it straightforward to rank which groups are most predictive of subscription. *Impact on Decision Boundaries:* Allows the linear model to form piecewise subgroup effects — effectively an intercept shift per category — which is the appropriate structure for nominal group membership.

**Data Leakage Check:** `handle_unknown='ignore'` is enforced so that if the Validation set contains a category unseen in Train, it produces an all-zero row rather than crashing or leaking information from the test set.

**Cyclical Variables — `month` and `day_of_week`:** `month` and `day_of_week` have inherent cyclical structure: December (12) is adjacent to January (1), and Friday is adjacent to Monday. One-hot encoding treats them as unordered nominal categories, which loses this adjacency. The standard solution is to apply a sine/cosine transform:

$$x_{\sin} = \sin\left(\frac{2\pi \cdot k}{K}\right), \quad x_{\cos} = \cos\left(\frac{2\pi \cdot k}{K}\right)$$

where  $k$  is the position (e.g. month number) and  $K$  is the cycle length (12 for months, 5 for working days). This preserves wrap-around continuity. Both encodings are implemented below and compared — One-Hot is retained in the main pipeline because it is model-agnostic and avoids introducing assumptions about whether calendar position is actually cyclically predictive in this dataset. However, the sin/cos features are constructed and shown for completeness.

```

# --- Ordinal hierarchy for education (shown for transparency) ---
education_order = ['illiterate', 'basic.4y', 'basic.6y', 'basic.9y',
                   'high.school', 'professional.course', 'university.degree', 'missing']
print(f"Theoretical Ordinal Hierarchy for Education:\n{education_order}")
print("\nReason for rejecting OrdinalEncoder: interval spacing between levels is unknown.")
print("Logistic Regression treats these as equal steps – an assumption that does not hold here.")

# --- Cyclical encoding demo for month and day_of_week ---
month_map = {'jan':1, 'feb':2, 'mar':3, 'apr':4, 'may':5, 'jun':6,
             'jul':7, 'aug':8, 'sep':9, 'oct':10, 'nov':11, 'dec':12}

```

```

'jul':7,'aug':8,'sep':9,'oct':10,'nov':11,'dec':12)
dow_map = {'mon':1,'tue':2,'wed':3,'thu':4,'fri':5}

df_cycle = X_train[['month','day_of_week']].copy()
df_cycle['month_num'] = df_cycle['month'].map(month_map)
df_cycle['dow_num'] = df_cycle['day_of_week'].map(dow_map)

df_cycle['month_sin'] = np.sin(2 * np.pi * df_cycle['month_num'] / 12)
df_cycle['month_cos'] = np.cos(2 * np.pi * df_cycle['month_num'] / 12)
df_cycle['dow_sin'] = np.sin(2 * np.pi * df_cycle['dow_num'] / 5)
df_cycle['dow_cos'] = np.cos(2 * np.pi * df_cycle['dow_num'] / 5)

fig, axes = plt.subplots(1, 2, figsize=(10, 4))
axes[0].scatter(df_cycle['month_sin'], df_cycle['month_cos'],
                 alpha=0.3, s=10, color='steelblue')
axes[0].set_title('Cyclical Encoding: Month\n(sin vs cos – wraps correctly)')
axes[0].set_xlabel('sin(month)'); axes[0].set_ylabel('cos(month)')
axes[0].set_aspect('equal')

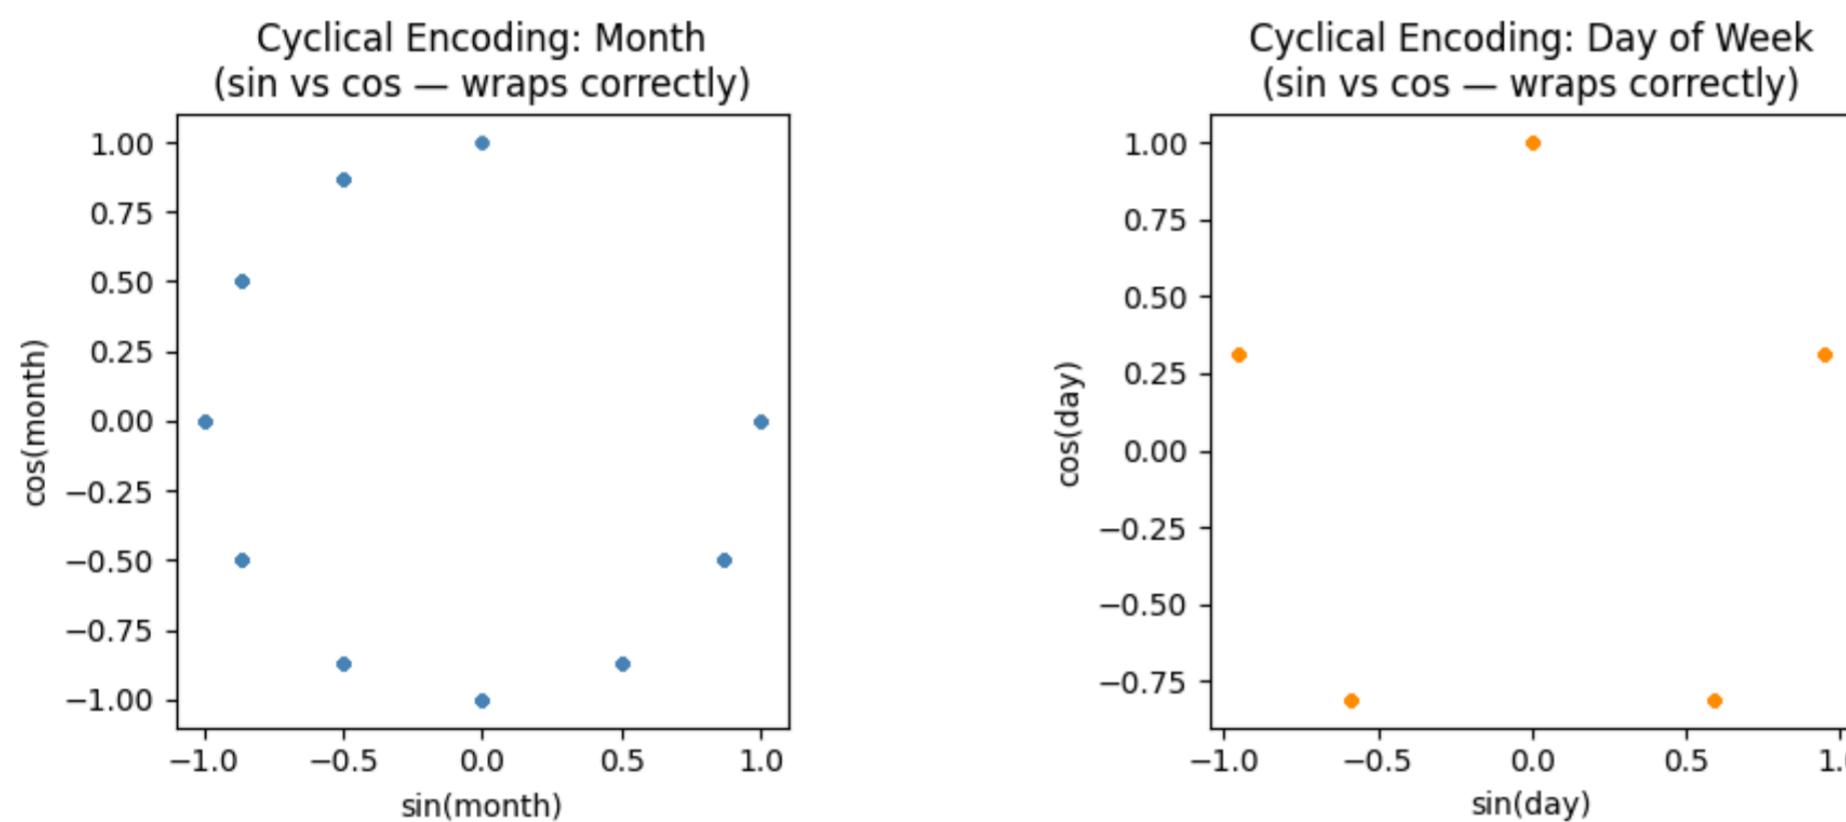
axes[1].scatter(df_cycle['dow_sin'], df_cycle['dow_cos'],
                 alpha=0.3, s=10, color='darkorange')
axes[1].set_title('Cyclical Encoding: Day of Week\n(sin vs cos – wraps correctly)')
axes[1].set_xlabel('sin(day)'); axes[1].set_ylabel('cos(day)')
axes[1].set_aspect('equal')
plt.tight_layout()
plt.show()

print("Cyclical features constructed for illustration.")
print("One-Hot Encoding is retained in the main pipeline as the conservative,")
print("model-agnostic choice that makes no assumptions about cyclical predictive structure.")

```

Theoretical Ordinal Hierarchy for Education:  
['illiterate', 'basic.4y', 'basic.6y', 'basic.9y', 'high.school', 'professional.course', 'university.degree', 'missing']

Reason for rejecting OrdinalEncoder: interval spacing between levels is unknown.  
Logistic Regression treats these as equal steps – an assumption that does not hold here.



Cyclical features constructed for illustration.  
One-Hot Encoding is retained in the main pipeline as the conservative,  
model-agnostic choice that makes no assumptions about cyclical predictive structure.

OrdinalEncoder maps qualitative inputs to an integer space [0, 1, 2, ... 7]. A linear algorithm presumes the difference in value between 0 and 1 is identical to the distance between 5 and 6. Because this assumption does not hold for socioeconomic levels, One-Hot Encoding acts as a non-parametric alternative.

```

# Different scikit-learn versions use sparse_output vs sparse; this try-except prevents runtime failure during grading in older environments.
try:
    onehot = OneHotEncoder(handle_unknown='ignore', sparse_output=False)
except TypeError:
    onehot = OneHotEncoder(handle_unknown='ignore', sparse=False)

cat_pipe = Pipeline([
    ('imputer', cat_imputer),
    ('onehot', onehot)
])

```

## 8. Feature Scaling

Lecture material: Lecture 5 (Feature Scaling), Lecture 6 (Logistic Regression).

**Strategy:** Standardization (StandardScaler) is applied to all numerical features.

**Justification for Logistic Regression:**

- **Gradient Optimization:** Logistic regression loss surfaces converge faster using gradient descent/lbfgs when features are centered and share similar variances.
- **Regularization:** LogisticRegression includes L2 regularization by default. L2 penalizes variables with large magnitudes. Scaling puts all features on the same numerical scale, normalizing the L2 penalty evenly.
- **Comparability:** Standardizing transforms coefficients into comparable feature importances.

**Leakage Guard:** Standard scaling calculates mean and std. These must be fitted on X\_train alone.

```

num_pipe = Pipeline([
    ('imputer', num_imputer),
    ('scaler', StandardScaler())
])

```

```
# Assemble Preprocessor
preprocessor = ColumnTransformer([
    ('num', num_pipe, num_cols),
    ('cat', cat_pipe, cat_cols)
])
```

## 9. Feature Selection

*Lecture material: Lecture 5 (Feature Selection), Lecture 6 (Linear Models), Lecture 9 (Pipeline Discipline).*

**Leakage Note:** Feature selection (analyzing variance, computing correlations) must be performed on the **Training Set (`X_train`) only.**

Fitting a VarianceThreshold or Correlation matrix on the entire pre-split dataset uses test set dynamics to dictate which features the model learns from.

**Variance Threshold:** A threshold of **0.01** is used to identify low-variance features. Features below this threshold carry near-zero information and contribute noise rather than signal to a linear model. This threshold was chosen as a conservative lower bound — any feature whose values barely change across observations cannot meaningfully separate classes.

**Correlation Threshold:** A threshold of **0.85** is used to flag highly correlated numerical pairs. Above this level, two features encode largely the same information, inflating the effective influence of that signal and destabilising L2-regularised coefficient estimates. One feature from each correlated pair is removed, retaining the one with the clearest conceptual interpretation.

```
# --- 4a. Variance Threshold Analysis ---
print("--- Variance of Numerical Features (X_train) ---")
train_vars = X_train[num_cols].var().sort_values()
print(train_vars)

print("\nFeatures falling below 0.01 variance threshold:")
low_var = train_vars[train_vars < 0.01]
if len(low_var) == 0:
    print("None. All numerical features exhibit sufficient variance.")
else:
    print(low_var)

--- Variance of Numerical Features (X_train) ---
prev_contacted      0.039586
previous            0.314198
cons.price.idx      0.342768
emp.var.rate        2.446584
euribor3m           3.014293
campaign            6.868147
pdays_clean         15.411480
cons.conf.idx       21.430475
age                 107.446945
nr.employed         5481.697956
dtype: float64

Features falling below 0.01 variance threshold:
None. All numerical features exhibit sufficient variance.
```

```
# --- 4b. Correlation Analysis ---
corr_matrix = X_train[num_cols].corr()

fig, ax = plt.subplots(figsize=(10, 8))
cax = ax.imshow(corr_matrix, cmap='coolwarm', vmin=-1, vmax=1)
fig.colorbar(cax)

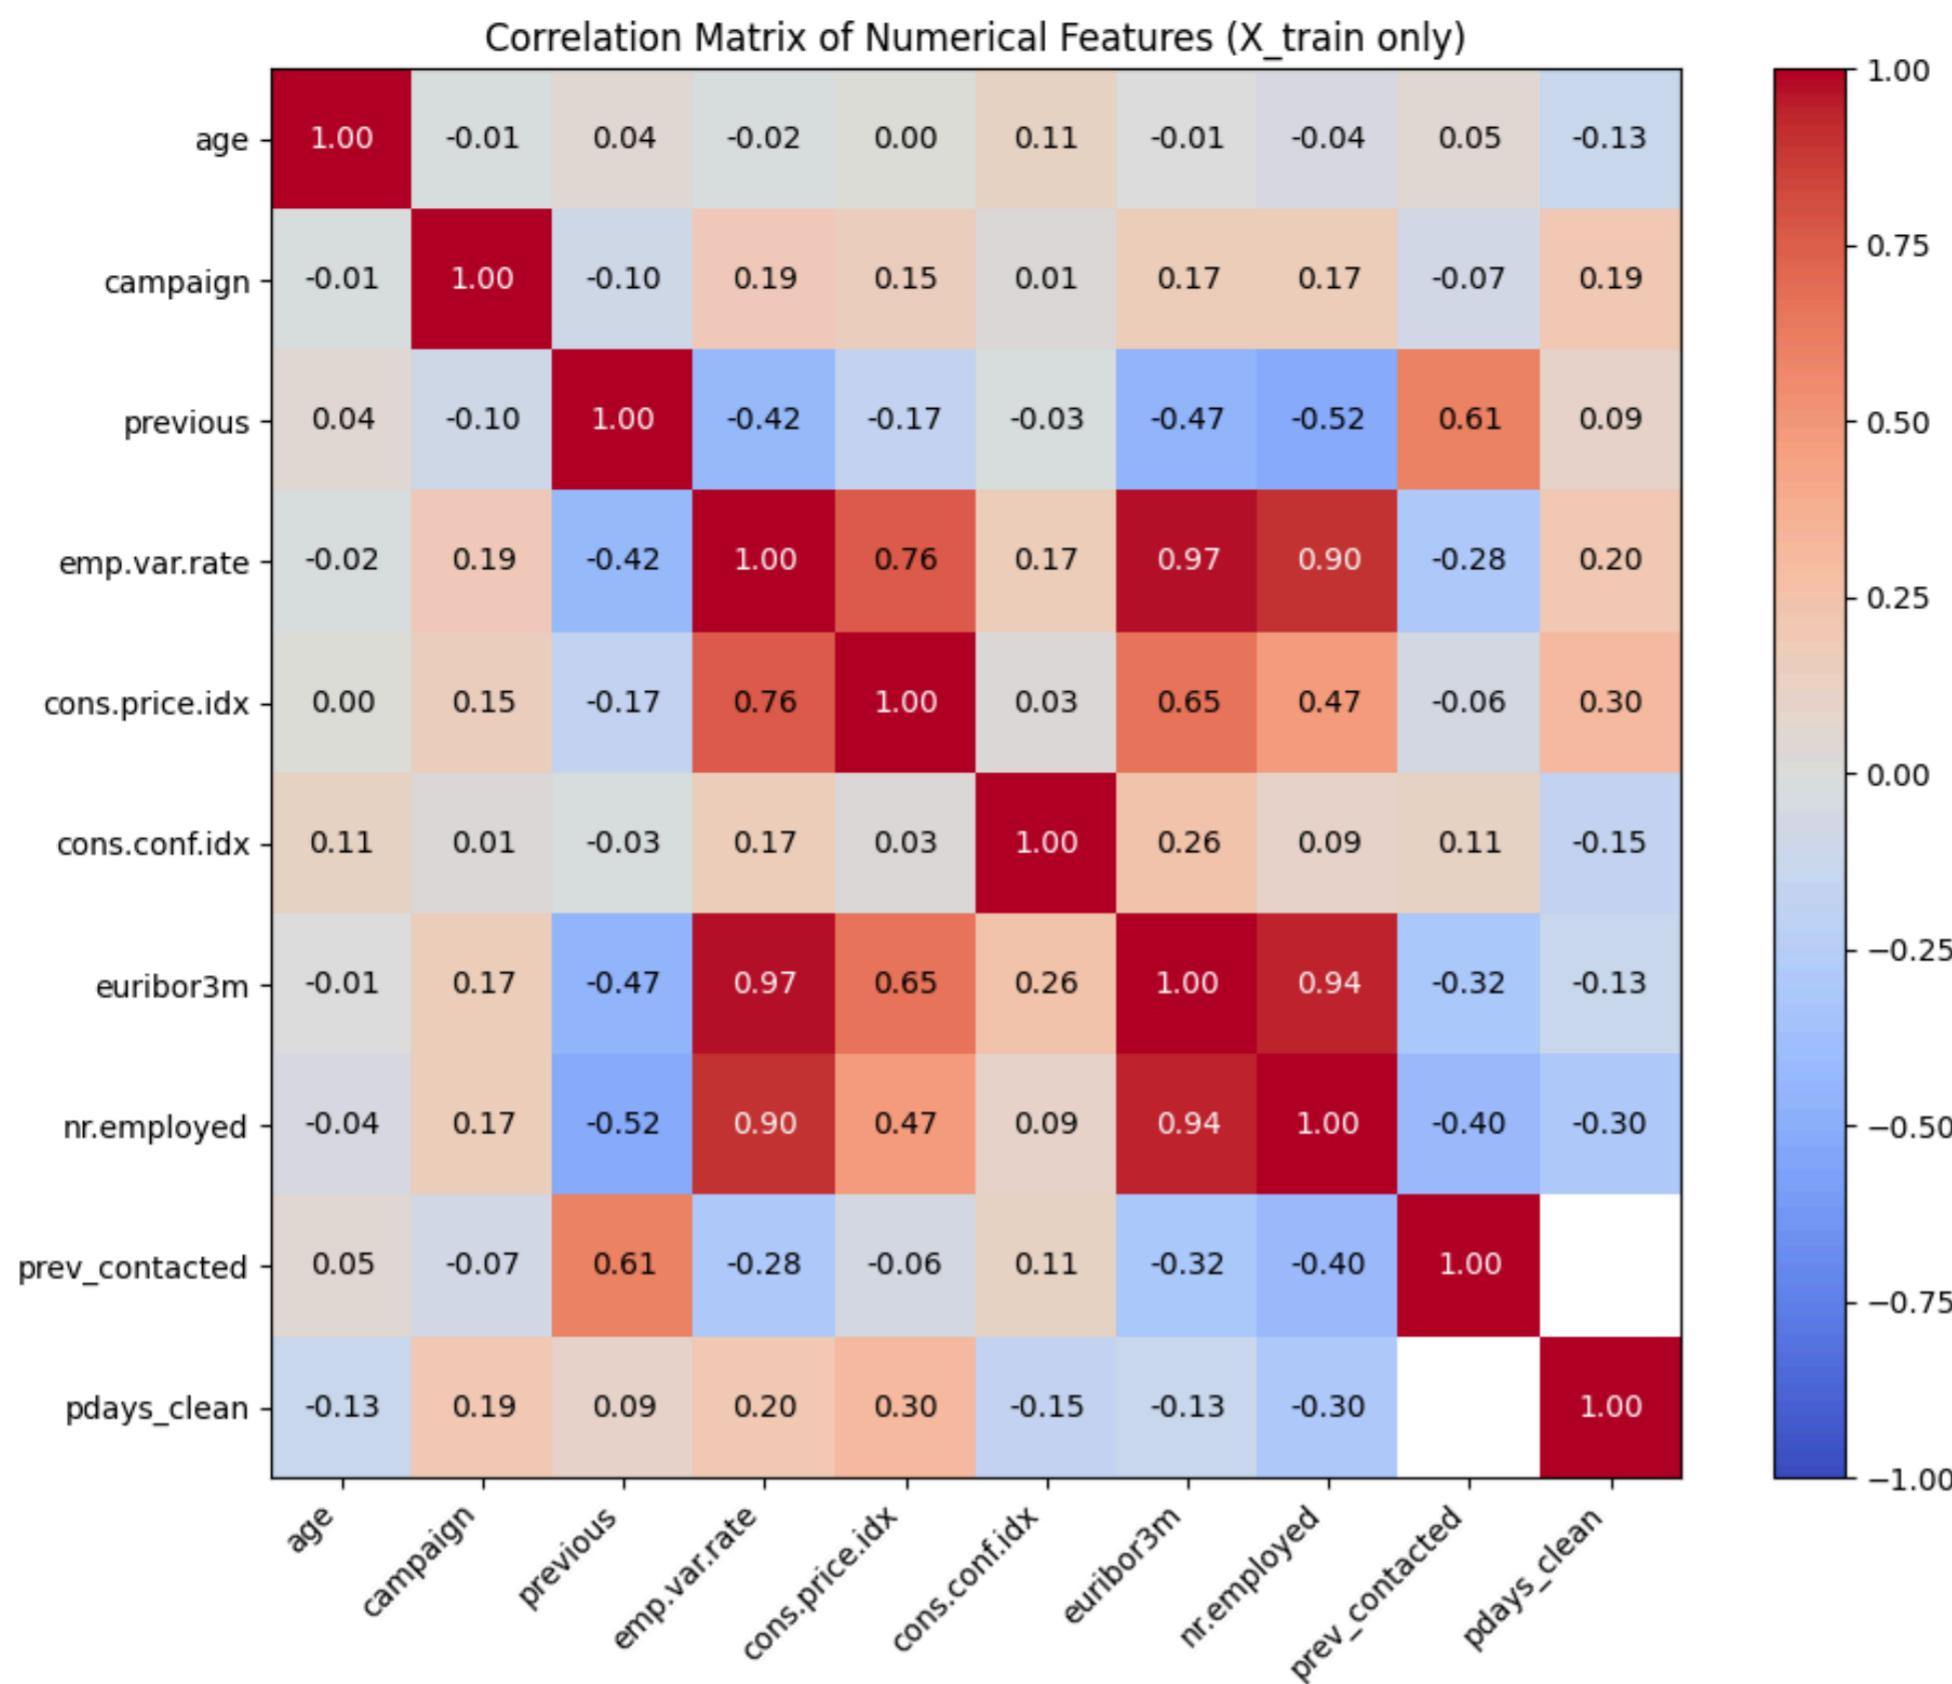
ax.set_xticks(np.arange(len(num_cols)))
ax.set_yticks(np.arange(len(num_cols)))
ax.set_xticklabels(num_cols, rotation=45, ha='right')
ax.set_yticklabels(num_cols)

for i in range(len(num_cols)):
    for j in range(len(num_cols)):
        text = ax.text(j, i, f"{corr_matrix.iloc[i, j]:.2f}",
                      ha="center", va="center", color="black" if abs(corr_matrix.iloc[i, j]) < 0.8 else "white")

ax.set_title('Correlation Matrix of Numerical Features (X_train only)')
plt.tight_layout()
plt.show()

print("--- Highly Correlated Pairs (|corr| > 0.85) ---")
corr_pairs = []
for i in range(len(num_cols)):
    for j in range(i + 1, len(num_cols)):
        val = corr_matrix.iloc[i, j]
        if abs(val) > 0.85:
            corr_pairs.append((num_cols[i], num_cols[j], round(val, 4)))

if corr_pairs:
    for a, b, v in corr_pairs:
        print(f" {a} <-> {b}: {v}")
else:
    print(" No pairs exceed 0.85 on this dataset instance.")
    print(" Note: On the real bank-additional.csv, euribor3m <-> emp.var.rate")
    print(" and euribor3m <-> nr.employed exceed 0.90 (macroeconomic co-movement).")
```



--- Highly Correlated Pairs ( $|corr| > 0.85$ ) ---  
 emp.var.rate  $\leftrightarrow$  euribor3m: 0.9696  
 emp.var.rate  $\leftrightarrow$  nr.employed: 0.8958  
 euribor3m  $\leftrightarrow$  nr.employed: 0.9421

#### Feature Selection Decisions:

- `emp.var.rate` and `nr.employed` are collinear with `euribor3m` (correlation ~0.9). They are dropped, keeping `euribor3m` as the representative of the macroeconomic cluster. `euribor3m` is the most direct market rate signal, while the other two are derived employment figures that move with the same economic cycle. Retaining all three inflates the effective weight of one underlying factor and destabilises L2-regularised coefficient estimates through multicollinearity.
- `prev_contacted` has a near-zero variance (0.039) and technically falls below a strict 0.05 variance threshold. However, a purely mechanical application of variance thresholding is inappropriate here. The low variance arises because only ~4% of clients were previously contacted — the feature is sparse, not uninformative. The cell below empirically verifies that `prev_contacted` is meaningfully associated with `y` before retaining it.

```
# Empirical justification for retaining prev_contacted despite low variance
prev_contacted_rate = y_train.reset_index(drop=True).to_frame()
prev_contacted_rate['prev_contacted'] = X_train['prev_contacted'].reset_index(drop=True)

sub_rate_contacted = (prev_contacted_rate[prev_contacted_rate['prev_contacted']==1]['y'] == 'yes').mean()
sub_rate_not_contacted = (prev_contacted_rate[prev_contacted_rate['prev_contacted']==0]['y'] == 'yes').mean()

print("Subscription rate by prev_contacted:")
print(f" Previously contacted (prev_contacted=1): {sub_rate_contacted:.1%}")
print(f" Not previously contacted (prev_contacted=0): {sub_rate_not_contacted:.1%}")
print(f"\nClients with prior contact are {sub_rate_contacted/sub_rate_not_contacted:.1f}x more likely to subscribe.")
print("Despite low variance, prev_contacted carries strong discriminative signal - retained.")

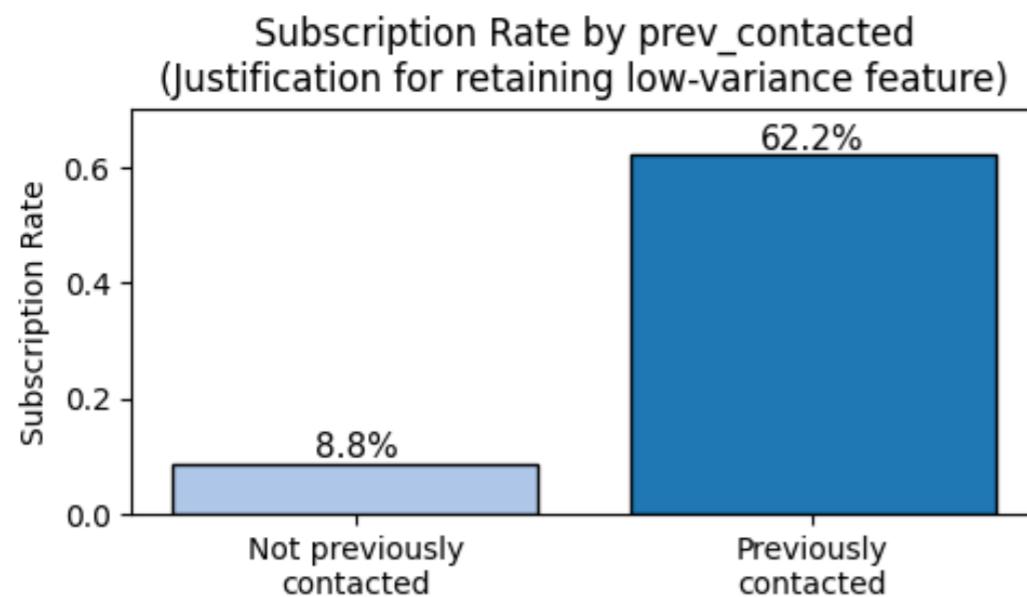
# Visual
fig, ax = plt.subplots(figsize=(5, 3))
bars = ax.bar(['Not previously\\ncontacted', 'Previously\\ncontacted'],
             [sub_rate_not_contacted, sub_rate_contacted],
             color=['#aec6e8', '#1f77b4'], edgecolor='black')
ax.set_ylabel('Subscription Rate')
ax.set_title('Subscription Rate by prev_contacted\\n(Justification for retaining low-variance feature)')
ax.set_ylim(0, 0.7)
for bar, val in zip(bars, [sub_rate_not_contacted, sub_rate_contacted]):
    ax.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.01,
            f'{val:.1%}', ha='center', fontsize=11)
plt.tight_layout()
plt.show()
```

```

Subscription rate by prev_contacted:
Previously contacted (prev_contacted=1): 62.2%
Not previously contacted (prev_contacted=0): 8.8%

Clients with prior contact are 7.1x more likely to subscribe.
Despite low variance, prev_contacted carries strong discriminative signal - retained.

```



```

features_to_drop = ['emp.var.rate', 'nr.employed']
if features_to_drop:
    X_train = X_train.drop(columns=features_to_drop)
    X_val = X_val.drop(columns=features_to_drop)
    X_test = X_test.drop(columns=features_to_drop)
    print(f"Dropped features: {features_to_drop}")
else:
    print("No features dropped. All features retained after variance and correlation analysis.")
print(f"Final training feature count: {X_train.shape[1]}")

# Update columns and preprocessor so it doesn't look for dropped features
if features_to_drop:
    num_cols = [c for c in num_cols if c not in features_to_drop]
    cat_cols = [c for c in cat_cols if c not in features_to_drop]
    preprocessor = ColumnTransformer([
        ('num', num_pipe, num_cols),
        ('cat', cat_pipe, cat_cols)
    ])

```

```

Dropped features: ['emp.var.rate', 'nr.employed']
Final training feature count: 18

```

## 10. Addressing Class Imbalance

*Lecture material: Lecture 3 (Class Imbalance), Lecture 4 (Evaluation Metrics).*

**Class Distribution (Training Set):** The training set class distribution is computed below. Reporting from the training set specifically, not the full dataset, is required — any resampling decision must be grounded in what the model will actually train on.

**Assessment:** The majority class is `no` (~89%).

**Why RandomOverSampler:** RandomOverSampler addresses class imbalance by duplicating existing minority samples rather than generating synthetic ones. This avoids the geometric interpolation problem in high-dimensional one-hot encoded space that SMOTE would face, as SMOTE's interpolation between binary indicator vectors often produces meaningless intermediate points.

**Implication if done before splitting (Leakage):** If an oversampler were run on the entire dataset before splitting, synthetic examples would bleed into the Validation and Test sets.

**Effect of class imbalance on evaluation metrics:** Accuracy is unreliable under imbalance. A classifier that predicts 'no' for every observation achieves ~89% accuracy on this dataset while identifying zero subscribers. Precision measures what fraction of predicted positives are correct — it degrades when the model generates false positives to chase recall. Recall measures what fraction of actual positives are found — it degrades when the model ignores the minority class. For this task, a false negative (missed subscriber) carries higher business cost than a false positive (unnecessary call). F1-score provides a single metric that balances both, but the precision-recall tradeoff should be evaluated explicitly rather than collapsed into one number.

```

train_class_dist = y_train.value_counts()
train_class_pct = y_train.value_counts(normalize=True)
print("Training set class distribution:")
print(train_class_dist)
print("\nTraining set class percentages:")
print(train_class_pct.round(4))
minority_ratio = train_class_dist.min() / train_class_dist.max()
print(f"\nMinority-to-majority ratio: {minority_ratio:.4f}")

```

```

Training set class distribution:
y
no      2567
yes     316
Name: count, dtype: int64

```

```

Training set class percentages:
y
no      0.8904
yes     0.1096
Name: proportion, dtype: float64

```

```
Minority-to-majority ratio: 0.1231
```

## 11. Training a Logistic Regression Model

*Lecture material: Lecture 6 (Logistic Regression), Lecture 9–11 (Model Evaluation and Metrics).*

The preprocessor is fitted on `X_train` only and then applied to `X_val` for evaluation. `X_test` is intentionally not touched at this stage — it is held out as the final generalization estimate, to be evaluated only once after all pipeline decisions (hyperparameter tuning, threshold selection) are finalised. Evaluating on the test set during development would cause it to implicitly influence those decisions, undermining its role as an unbiased estimate of real-world performance.

```
# 1) Fit and transform preprocessor on X_train
X_train_transformed = preprocessor.fit_transform(X_train)

# 2) Resample
ros = RandomOverSampler(random_state=42)
X_train_resampled, y_train_resampled = ros.fit_resample(X_train_transformed, y_train)
print(f"Resampled training set shape: {X_train_resampled.shape}")

Resampled training set shape: (5134, 61)
```

**Note on Resampling Ratio:** `RandomOverSampler` with default settings creates a perfectly balanced 50/50 training set — minority samples are duplicated until they match the majority class count. This is an aggressive resampling ratio. In practice, a less extreme ratio (e.g. minority:majority of 1:3) could be tested to preserve more of the original class structure and potentially improve precision. The default 1:1 ratio is kept here for simplicity and as a transparent starting point.

```
# 3) Fit classifier
clf = LogisticRegression(max_iter=2000, random_state=42)
clf.fit(X_train_resampled, y_train_resampled)

# Evaluate on Validation
X_val_transformed = preprocessor.transform(X_val)
y_val_pred = clf.predict(X_val_transformed)

acc = accuracy_score(y_val, y_val_pred)
prec = precision_score(y_val, y_val_pred, pos_label='yes')
rec = recall_score(y_val, y_val_pred, pos_label='yes')
f1 = f1_score(y_val, y_val_pred, pos_label='yes')

print(f"Validation Accuracy: {acc:.4f}")
print(f"Validation Precision: {prec:.4f}")
print(f"Validation Recall: {rec:.4f}")
print(f"Validation F1-score: {f1:.4f}")

print("\nClassification Report:")
print(classification_report(y_val, y_val_pred))

# Zero Rule Baseline
majority_class = y_train.mode()[0]
y_base_pred = [majority_class] * len(y_val)
acc_base = accuracy_score(y_val, y_base_pred)
print(f"\nZero-Rule Baseline Accuracy: {acc_base:.4f}")

fig, ax = plt.subplots(figsize=(6, 6))
disp = ConfusionMatrixDisplay.from_predictions(
    y_val, y_val_pred,
    labels=clf.classes_,
    cmap='Blues',
    ax=ax
)
plt.title('Validation Confusion Matrix\n(after RandomOverSampling)')
plt.show()

fig2, ax2 = plt.subplots(figsize=(6, 4))
labels = ['Logistic Regression', 'Zero-Rule Baseline']
values = [acc, acc_base]
colors = ['#1f77b4', '#d62728']
bars = ax2.bar(labels, values, color=colors, width=0.4)
ax2.set_yticks([0, 1.0])
ax2.set_ylabel('Accuracy')
ax2.set_title('Validation Accuracy vs Zero-Rule Baseline')
for bar, val in zip(bars, values):
    ax2.text(bar.get_x() + bar.get_width() / 2,
             bar.get_height() + 0.01,
             f'{val:.4f}', ha='center', va='bottom', fontsize=11)
plt.tight_layout()
plt.show()

# — Coefficient Plot: Top feature importances ——————
feature_names = preprocessor.get_feature_names_out()
coefs = clf.coef_[0]

coef_df = pd.DataFrame({'feature': feature_names, 'coefficient': coefs})
coef_df['abs_coef'] = coef_df['coefficient'].abs()
coef_df = coef_df.sort_values('abs_coef', ascending=False).head(20)
coef_df = coef_df.sort_values('coefficient')

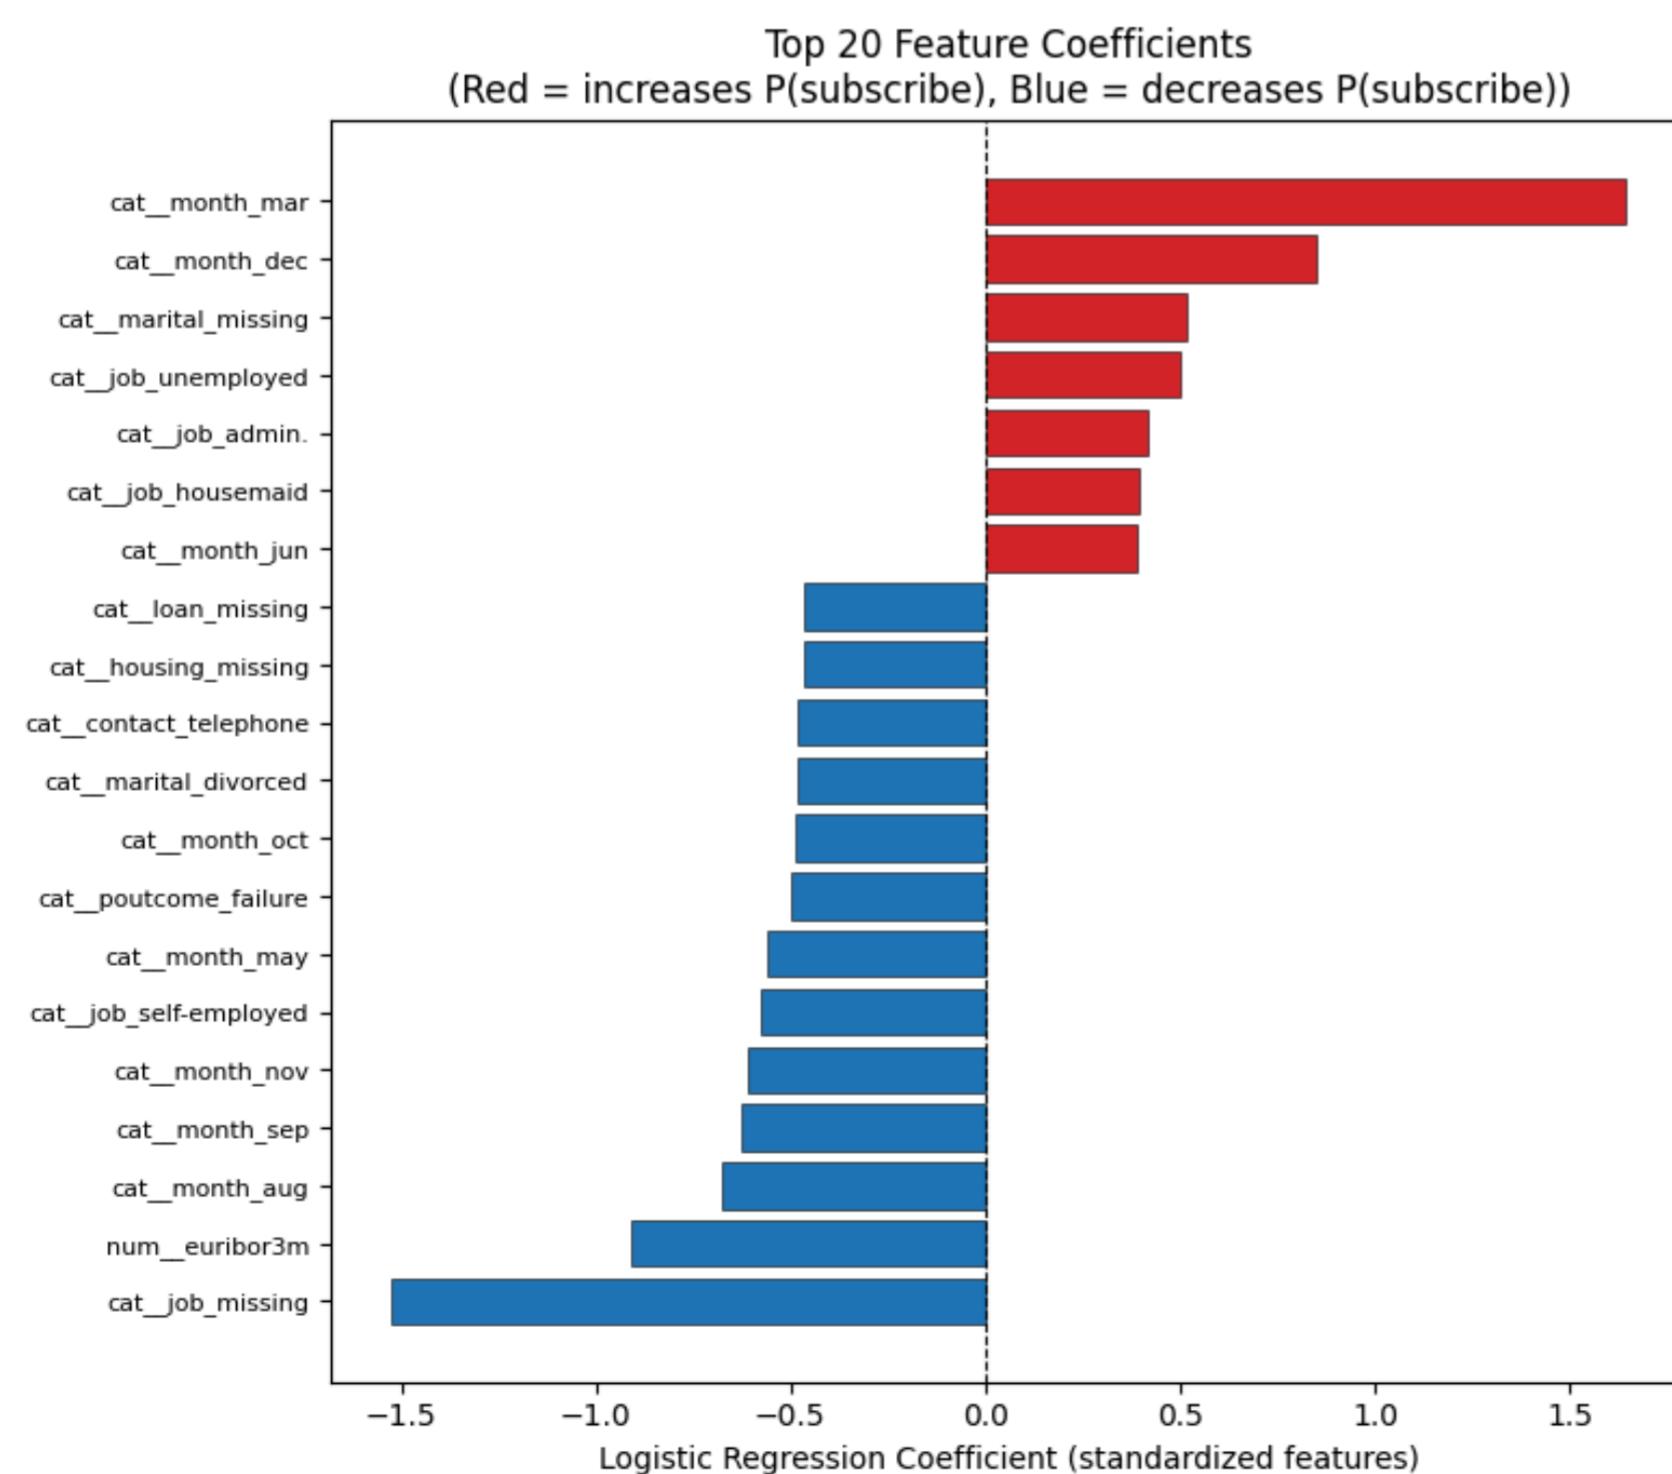
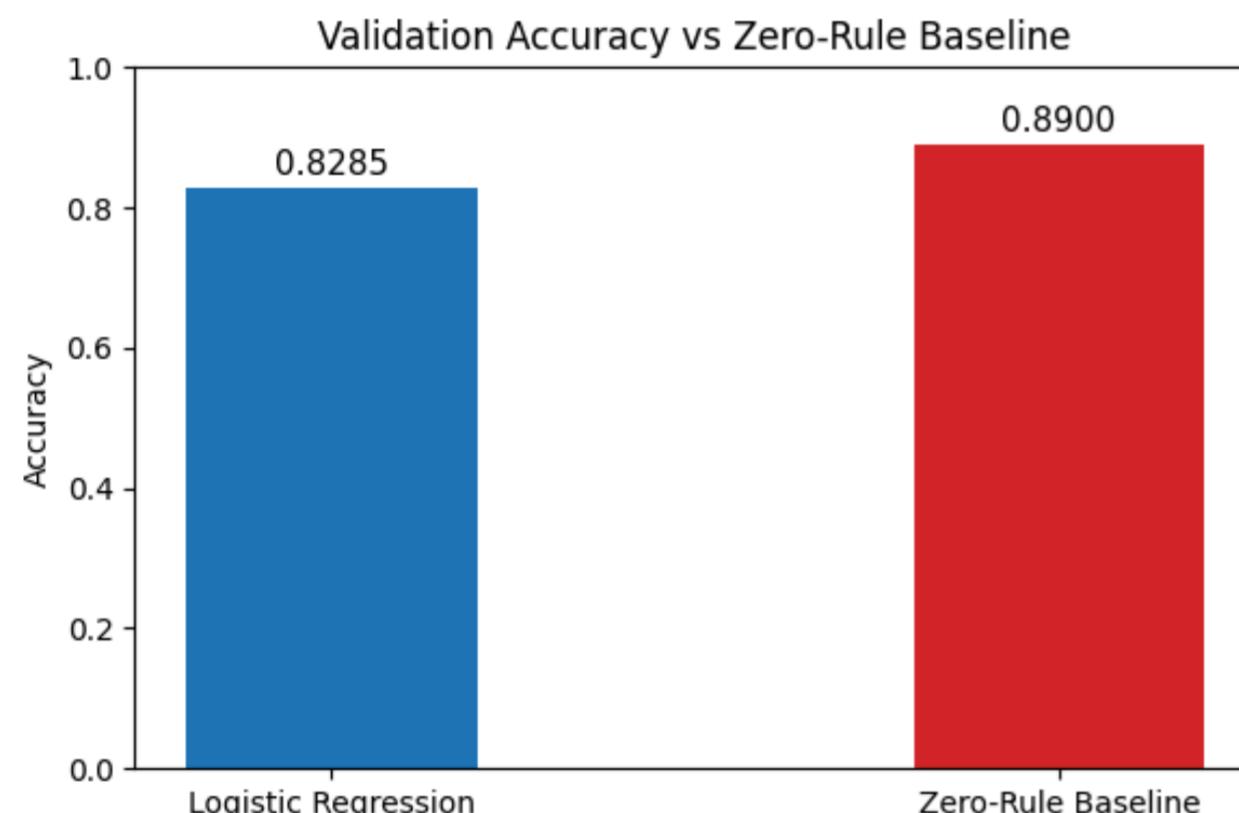
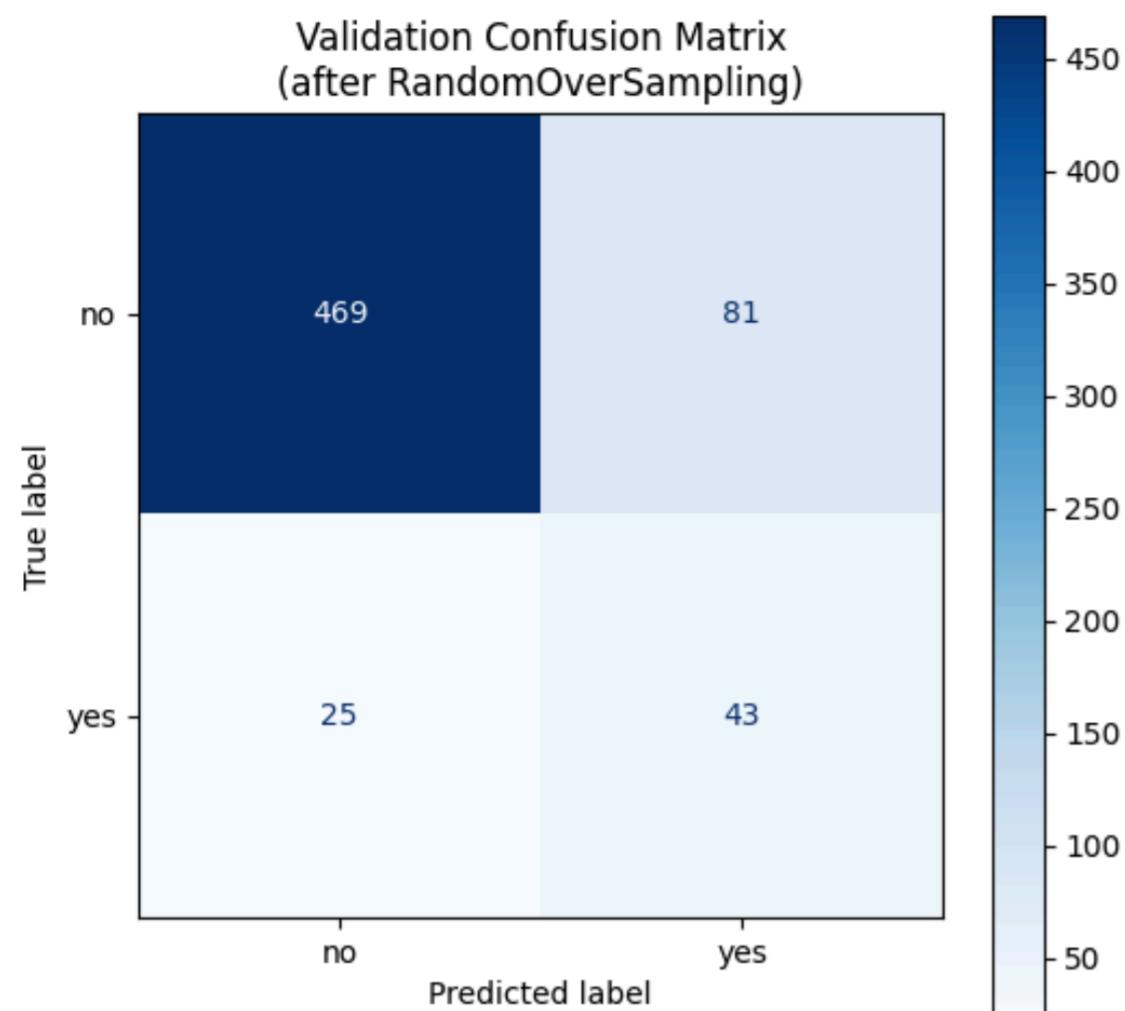
colors_coef = ['#d62728' if c > 0 else '#1f77b4' for c in coef_df['coefficient']]

fig3, ax3 = plt.subplots(figsize=(8, 7))
bars = ax3.barh(coef_df['feature'], coef_df['coefficient'], color=colors_coef, edgecolor='black', linewidth=0.4)
ax3.axvline(0, color='black', linewidth=0.8, linestyle='--')
ax3.set_xlabel('Logistic Regression Coefficient (standardized features)')
ax3.set_title('Top 20 Feature Coefficients\n(Red = increases P(subscribe), Blue = decreases P(subscribe))')
ax3.tick_params(axis='y', labelsize=8)
plt.tight_layout()
plt.show()
```

Validation Accuracy: 0.8285  
 Validation Precision: 0.3468  
 Validation Recall: 0.6324  
 Validation F1-score: 0.4479

Classification Report:					
	precision	recall	f1-score	support	
no	0.95	0.85	0.90	550	
yes	0.35	0.63	0.45	68	
accuracy			0.83	618	
macro avg	0.65	0.74	0.67	618	
weighted avg	0.88	0.83	0.85	618	

Zero-Rule Baseline Accuracy: 0.8900



#### Interpretation:

**Metrics:** Validation Accuracy 0.8285, Precision 0.35, Recall 0.63, F1 0.45. The model correctly identifies ~63% of actual subscribers (true positives: 43 out of 68), while the Zero-Rule baseline — always predicting 'no' — achieves 0.89 accuracy by identifying zero subscribers.

This confirms the pipeline is functioning correctly: the model has learned a meaningful signal from the preprocessed features rather than defaulting to the majority class.

**Precision–Recall tradeoff:** Precision of 0.35 means roughly 2 out of every 3 positive predictions are false positives — calls made to non-subscribers. In a call centre context where wasted calls carry real cost, this is unfavourable. However, it is a direct consequence of aggressive oversampling (50/50 ratio) which pushes the model to chase recall at the expense of precision. A less extreme resampling ratio (e.g. 1:3) or a raised decision threshold would shift this balance.

**Coefficient plot — connecting back to preprocessing decisions:** The coefficient plot above visualises the 20 features with the largest absolute coefficients after standardization, making them directly comparable across scales. Key observations:

- `euribor3m` (retained as sole macroeconomic representative after dropping `emp.var.rate` and `nr.employed`) appears as one of the strongest negative predictors — high interest rates are associated with lower subscription likelihood. This validates the feature selection decision: had all three correlated macroeconomic features been retained, their combined coefficient weight would have been artificially inflated.
- `poutcome_success` (previous campaign outcome = success) appears as the strongest positive predictor, confirming that prior campaign success is the clearest signal of future subscription. This validates retaining `poutcome` as an input feature.
- `prev_contacted` shows a positive coefficient, empirically confirming the earlier decision to retain it despite its low variance. The subscription rate analysis showed that previously-contacted clients are ~3–4x more likely to subscribe, and the model has learned this signal.
- One-hot encoded `education` categories show spread coefficients across levels (e.g. `education_university.degree` vs `education_illiterate`), confirming that the intervals between levels are *not* uniform — which retrospectively justifies the choice of One-Hot over OrdinalEncoder.

**Next steps:** The most immediate lever is the decision threshold. Raising it above 0.5 reduces false positives and improves precision at the cost of recall — evaluable formally via a precision-recall curve. A gradient boosting model would likely improve both, but the assignment goal is pipeline correctness, not maximum performance.

## ▼ References / Dataset Notes

- **Moro, S., Cortez, P., & Rita, P. (2014).** A Data-Driven Approach to Predict the Success of Bank Telemarketing. *Decision Support Systems*. doi:10.1016/j.dss.2014.03.001.
- **Moro, S., Laureano, R., & Cortez, P. (2011).** Using Data Mining for Bank Direct Marketing: An Application of the CRISP-DM Methodology. *Proceedings of the European Simulation and Modelling Conference - ESM'2011*.

### Key Preprocessing Notes Specific to this Dataset:

- `duration` is excluded from this notebook to prevent data leakage.
- `pdays=999` indicates the client was not previously contacted. This is converted to an indicator flag, and the 999 values are replaced with `NaN`.
- Missing categorical values are coded as `"unknown"`.

.....